


Final RC Part 3

VE370FA22 TA Jiajun Gu

N-way Set Associative Cache

Motivation: Reduce miss rate




- Each set contains n blocks
- A main memory block can use *any of the blocks* within the corresponding set
- $Set\ Index = (Block\ address) \% (number\ of\ sets\ in\ cache)$
- $Block\ Offset = word\ offset + byte\ offset$

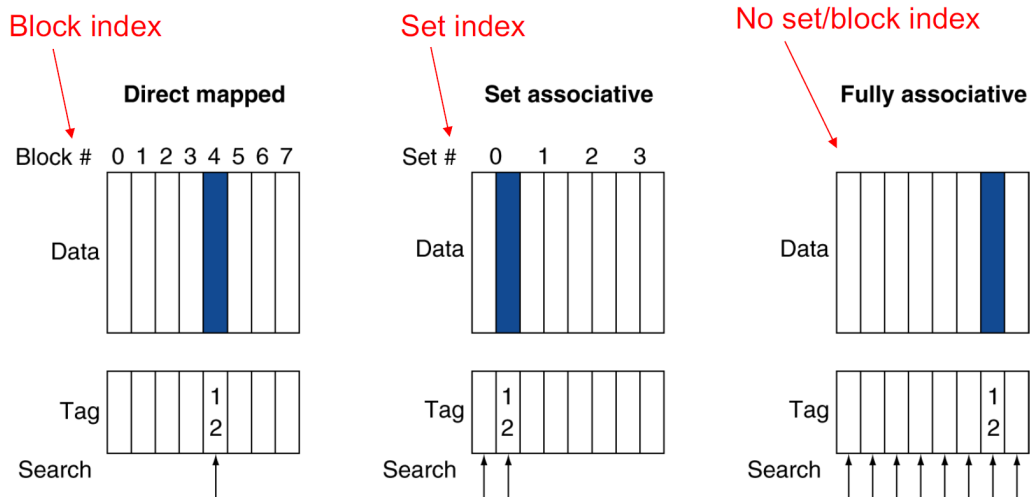
CAUTION

- n-way associative means each set contains n blocks, not the cache contains n sets!
- The basic element for operation is still **block** (compare with **page** in VM)

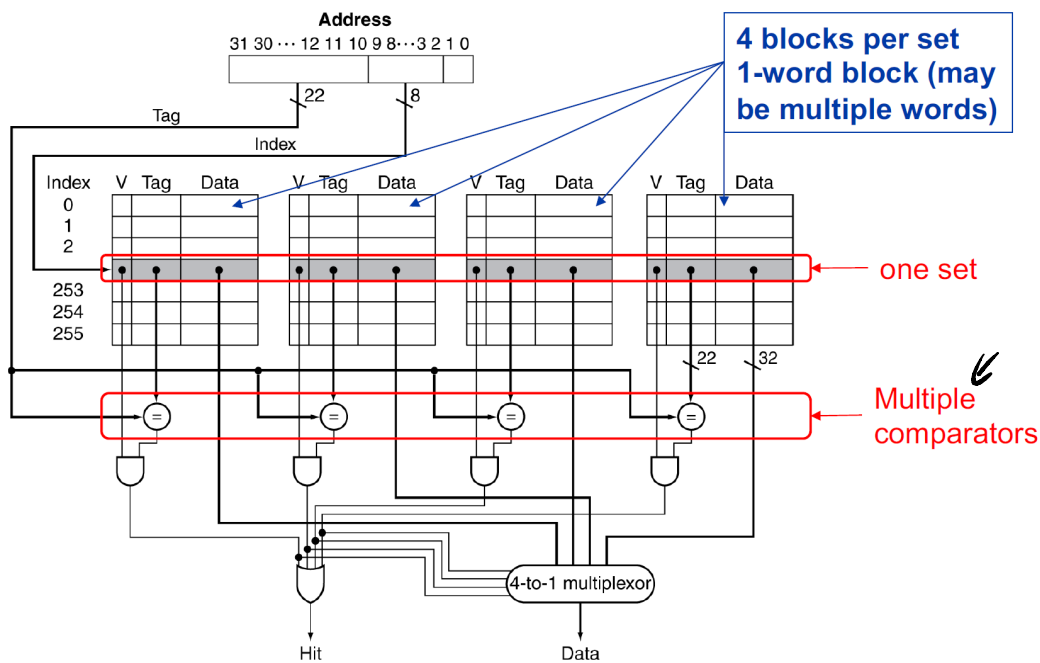
Special case



- $n = 1$: Directed mapped
- $n = total\ num\ of\ blocks$: Fully associative



Tradeoff: We may need to search n times in a set / use specific circuit structure (comparators) to locate a block in it. -> Extra hit time / hardware



Exercise:

- Suppose memory size is 64 words, use 4-way associative cache of 16 words size, 2 words per block.

- Memory

- memory word address: $\log_2(64) = 6$ bits

- memory byte address: $6 + 2 = 8$ bits

- Cache

- byte offset: 2 bits

- word offset: 1 bit

- set index: 1 bit

- number of blocks: $16/2 = 8$

- number of sets: $8/4 = 2$

- tag: $8 - 2 - 1 - 1 = 4$ bits

- Given a 4-way associative cache with 2K 8-word blocks, and a 32-bit byte address 0x810023FE, show the set index and tag for this byte address. ($2K=2^{11}$)

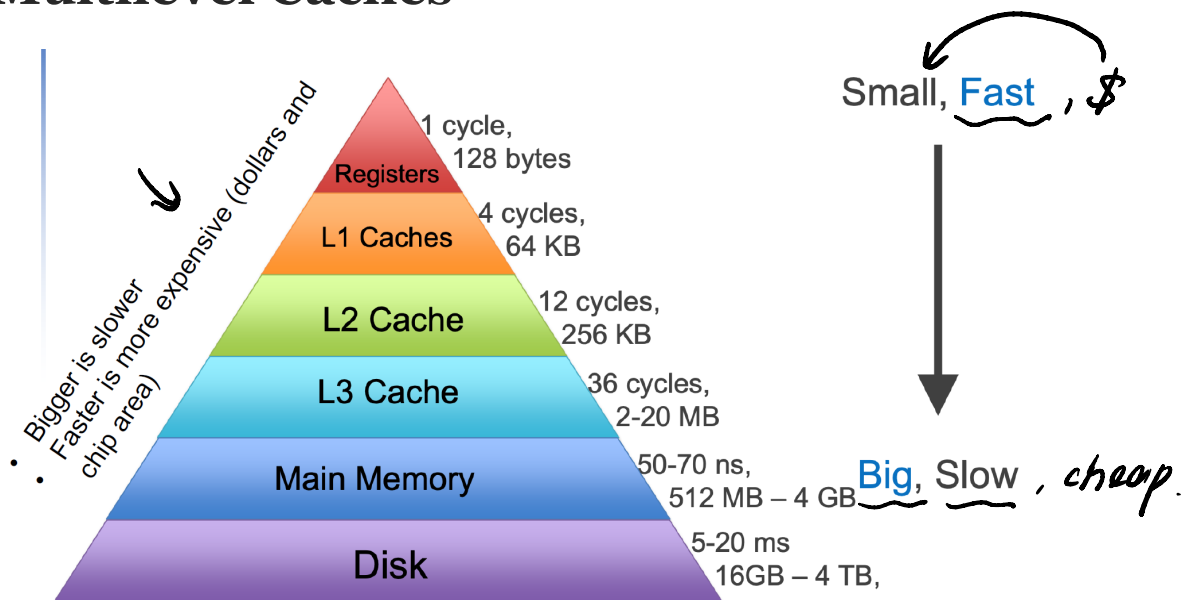
Tag	Set Index	Word Offset	Byte Offset
31:14	13:5	4:2	1:0

Replacement Policy

- Direct mapped: no other choices
- Set associative:
 - Prefer non-valid (empty) entry, if there is one
 - Otherwise, choose to replace a block in the set
- Policy:
 - **Least-recently used (LRU)** ←
 - Choose the one unused for the longest time
 - Simple for 2-way, manageable for 4-way, too hard beyond that
 - Random ←
 - Gives approximately the same performance as LRU **for high associativity**

Example: (Check by yourself) T11 p35-41

Multilevel Caches



- Primary (L-1) cache attached to CPU
 - Small, but fast
 - Focus on **minimal hit time** because miss penalty is smaller
 - Smaller block size *↳ Total cache size is smaller.*
 - Reduce search (hit) time
 - Reduce miss penalty (less time to fetch)

- Level-2 (secondary) cache services misses from primary cache
 - Larger, slower, but still faster than main memory
 - Focus on **low miss rate** to avoid main memory access
 - Higher associativity and block size

Memory Unit	Function (likely)	SRAM/DRAM
L1 cache	A cache for a cache	SRAM
L2 cache	A cache for main memory	SRAM
Main Memory	A cache for disks	DRAM
TLB	A cache for a page table	SRAM

Modification	Effect on miss rate	Negative Effect
① Increase cache size	Decrease <u>capacity misses</u>	Increase cost/access time (if use other materials/structures)
② Increase associativity	Decrease <u>conflict misses</u>	Increase complexity/hit time
③ Increase block size	Decrease <u>compulsory misses</u> (May increase conflict misses (e.g., cache size fixed))	Increase miss penalty

More on miss types: Types of Cache Misses - GeeksforGeeks

Performance Evaluation (multilevel cache):

CPI

If not specified, base CPI should contain the CPI used for hit.

$$\text{Miss penalty (main memory)} = \frac{\text{Main Memory Access Time}}{\text{Clock Cycle Time}}$$

With L-1 cache:

$$\text{Effective CPI} = \text{Base CPI} + \text{Miss penalty (main memory)} \times \text{Miss Rate}$$

With L-2 cache:

$$\text{Access Time (L-1 to L-2)} = \text{L-1 Miss penalty}$$

$$\text{Miss penalty (L-1 to L-2)} = \left\lceil \frac{\text{Access Time (L-1 to L-2)}}{\text{Clock Cycle Time}} \right\rceil$$

$$\begin{aligned} \text{Miss penalty (L-1 miss \&\& L-2 miss)} \\ = \text{Miss penalty (L-1 to L-2)} + \text{Miss penalty (main memory)} \end{aligned}$$

$$CPI = Base\ CPI$$

$$+ L-1\ Miss\ rate \times Miss\ penalty\ (L-1\ to\ L-2) \times L-2\ hit\ rate$$

$$+ L-1\ Miss\ rate \times Miss\ penalty\ (L-1\ miss\ \&\&\ L-2\ miss) \times L-2\ miss\ rate$$

Performance ratio = CPI (with L-1 cache) / CPI (with L-1 & L-2)

AMAT

$$AMAT\ (one\ level) = Hit\ Time + Miss\ Rate \times Miss\ Penalty$$

If not specified, hit time should also be spent when a miss occurs. (Think about this: why **Hit Time** is not multiplied by **Hit Rate**?)

$$AMAT\ (two\ levels) = L1\ Hit\ Time$$

$$+ L1\ Miss\ Rate \times (L1\ Miss\ Penalty + L2\ Miss\ Rate \times L2\ Miss\ Penalty)$$

- (20 points) Assume that main memory accesses take 70 ns and that memory accesses are 36% of all instructions. The following table shows parameters for a two-level cache memory.

	Size	Miss Rate	Hit Time
L1	16 KB	7.3%	1.18 ns
L2	1 MB	1.5%	5.34 ns

- What is the AMAT for the computer? (10 points)

Answer:

$$\begin{aligned} AMAT &= L1\ hit\ time + L1\ miss\ rate * L1\ miss\ penalty + L2\ miss\ rate * memory\ access\ time \\ &= 1.18 + 7.3\% * 98.5\% * 5.34 + 7.3\% * 1.5\% * (5.34 + 70) = 1.18 + 0.384 + 0.083 = 1.647\ ns \end{aligned}$$

- Assuming the L1 hit time determines the cycle times and a base CPI is 1.0 without any memory stalls, what is the total CPI? (10 points)

Answer: assuming only data cache is considered

$$L1\ miss\ penalty = 5.34 / 1.18 = 5\ cycles$$

$$L2\ miss\ penalty = 70 / 1.18 = 60\ cycles$$

$$Total\ CPI = base\ CPI + 36\% * (7.3\% * 98.5\% * 5 + 7.3\% * 1.5\% * (5 + 60)) = 1.15$$

Other notes (may not be useful)

What does it mean for an architecture to be called a load/store architecture?

A. Load and Store instructions are supported by the ISA.

- B. Load and Store instructions can also perform arithmetic instructions on data in memory.
- C. Data must first be loaded into a register before it can be operated on.
- D. Every load must have an accompanying store at some later point in the program.

Reference

1. VE370 2021FA, Final RC III
2. ECE3700J 2022FA, Lecture Slides