

Topic 13

Memory Hierarchy **- Virtual Memory (2)**

A Problem

- Page table is located in main memory
- Address translation requires an extra main memory access
 - One to access the page table
 - Then the actual memory access

Fast Translation Using a TLB

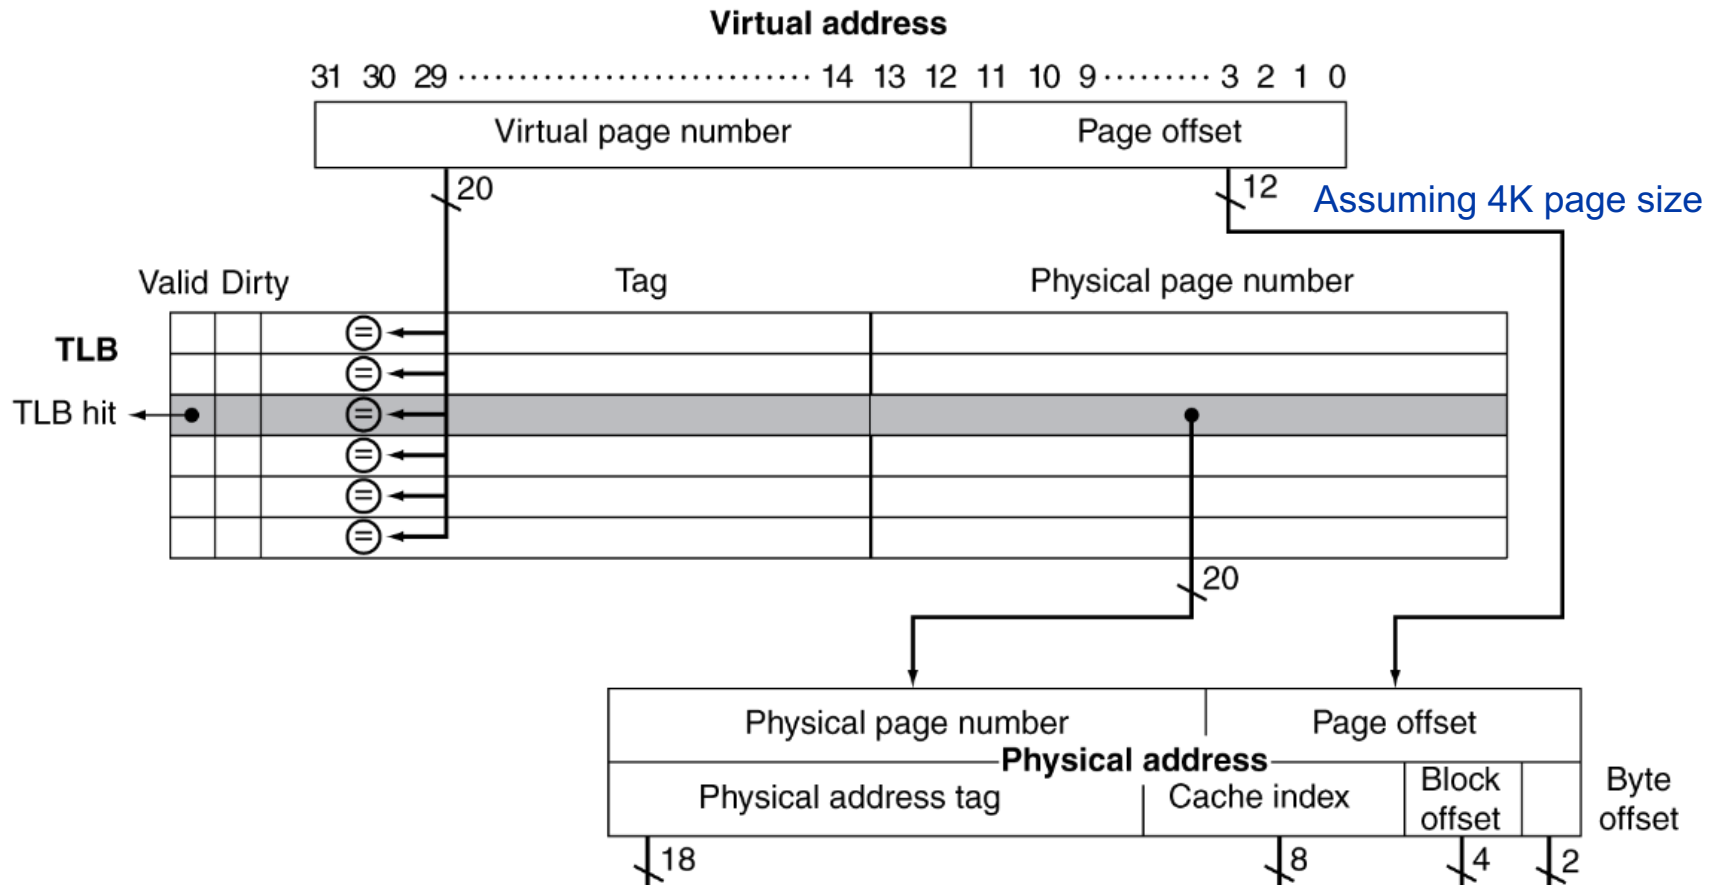
- Access to page tables has good locality
 - So use a fast **cache** of page tables within the CPU to store a subset of the page table translations
 - Called a **Translation Look-aside Buffer (TLB)**
 - The relationships are comparable:
 $(\text{TLB} \leftarrow \text{page table}) = (\text{cache} \leftarrow \text{main memory})$

TLB

- Load part of the Page Table in TLB
- In TLB, one translation per entry
- Typical: 16–512 entries, 0.5–1 cycle for hit, 10–100 cycles for miss, 0.01%–1% miss rate
- Status bits
 - Valid: empty or not
 - Dirty: target memory is changed
 - Reference: target memory is used
- Full Associativity
 - Lower miss rate
 - Small TLB, so access time not a major concern
- LRU or random replacement

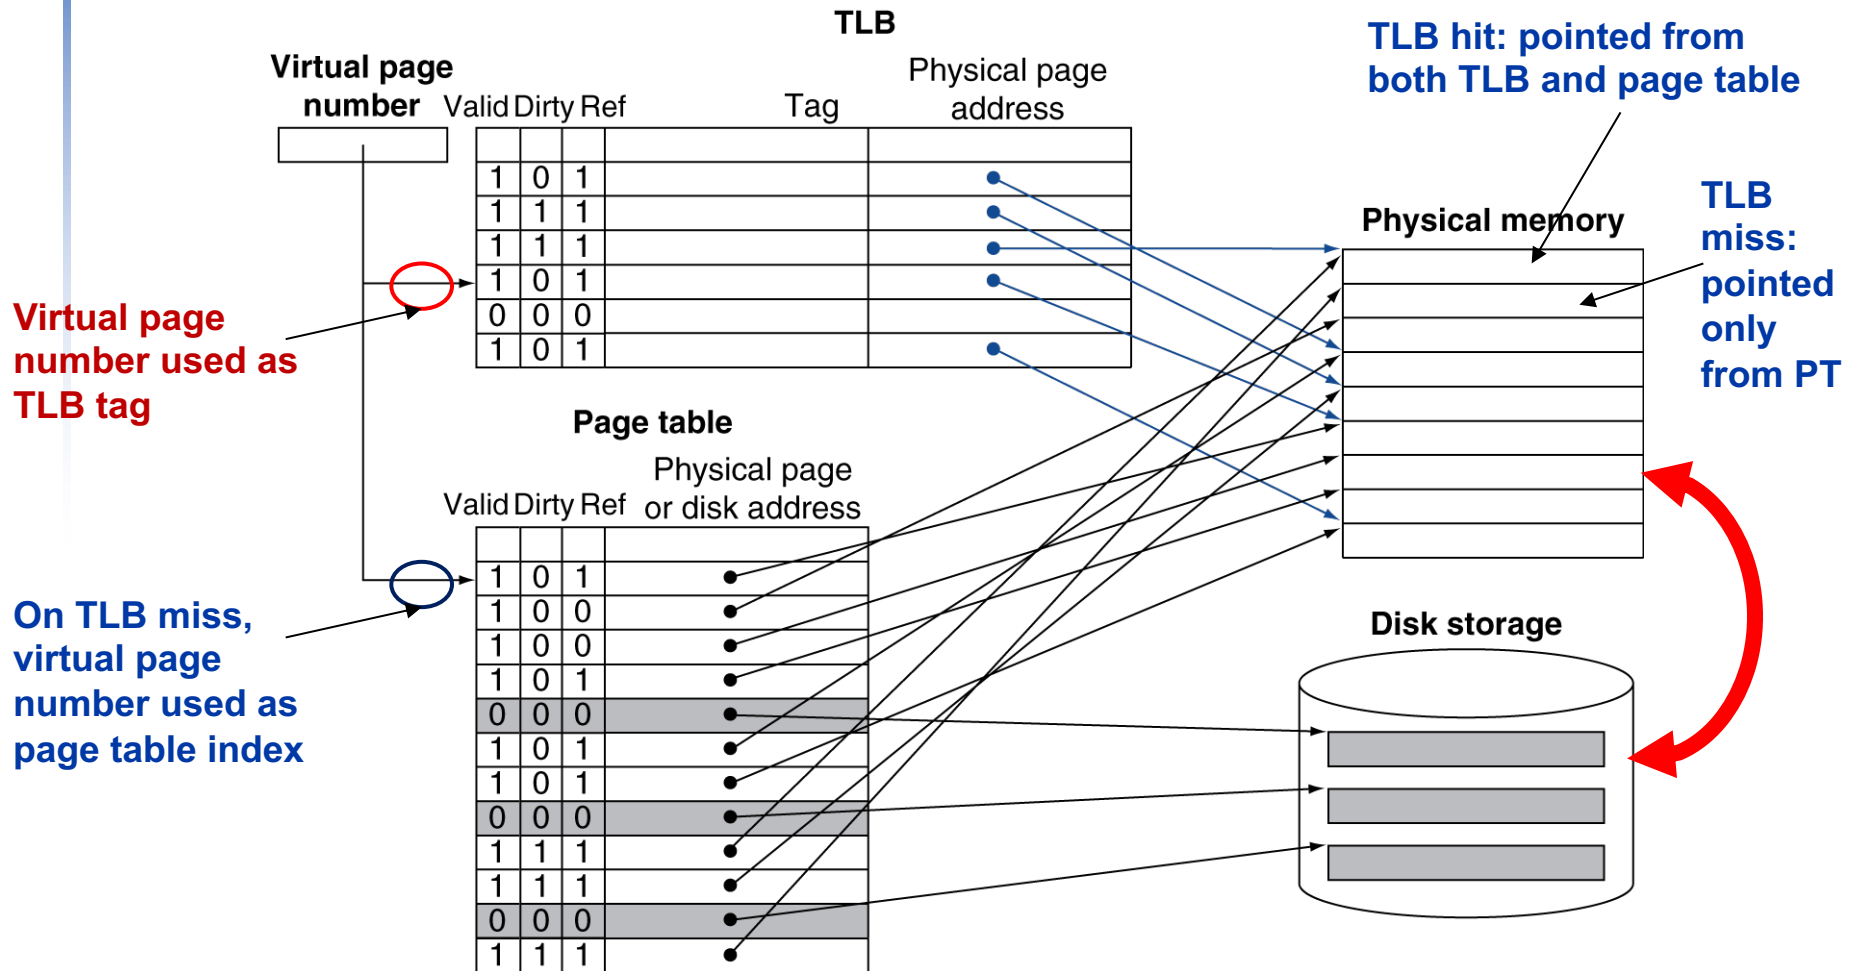
TLB Translation

Virtual page number is the Tag that is compared with all Tag fields of TLB because of full associativity of TLB (only 1 set)

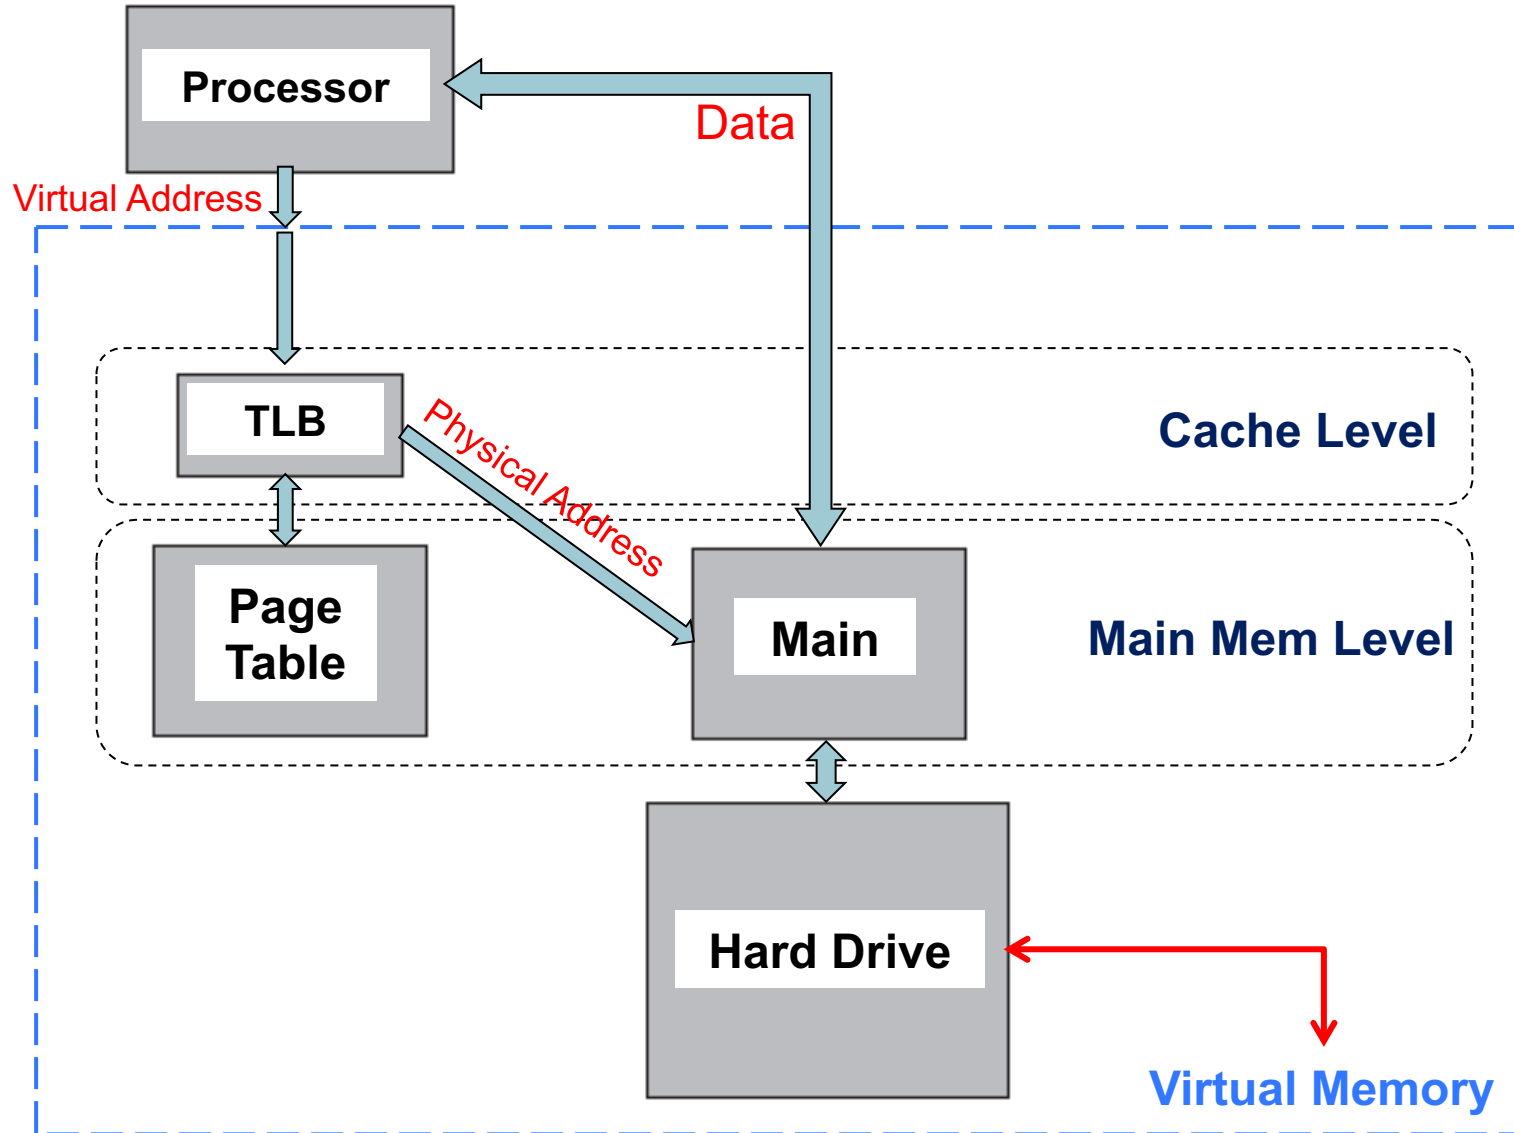


To access physical memory

Fast Translation Using a TLB



Access Mem through TLB



TLB Hit

- Virtual address found in TLB – hit
 - Provide physical address
 - Reference bit on
 - Dirty bit on if physical address used for write (might not be needed if a dirty bit is maintained in physical memory)

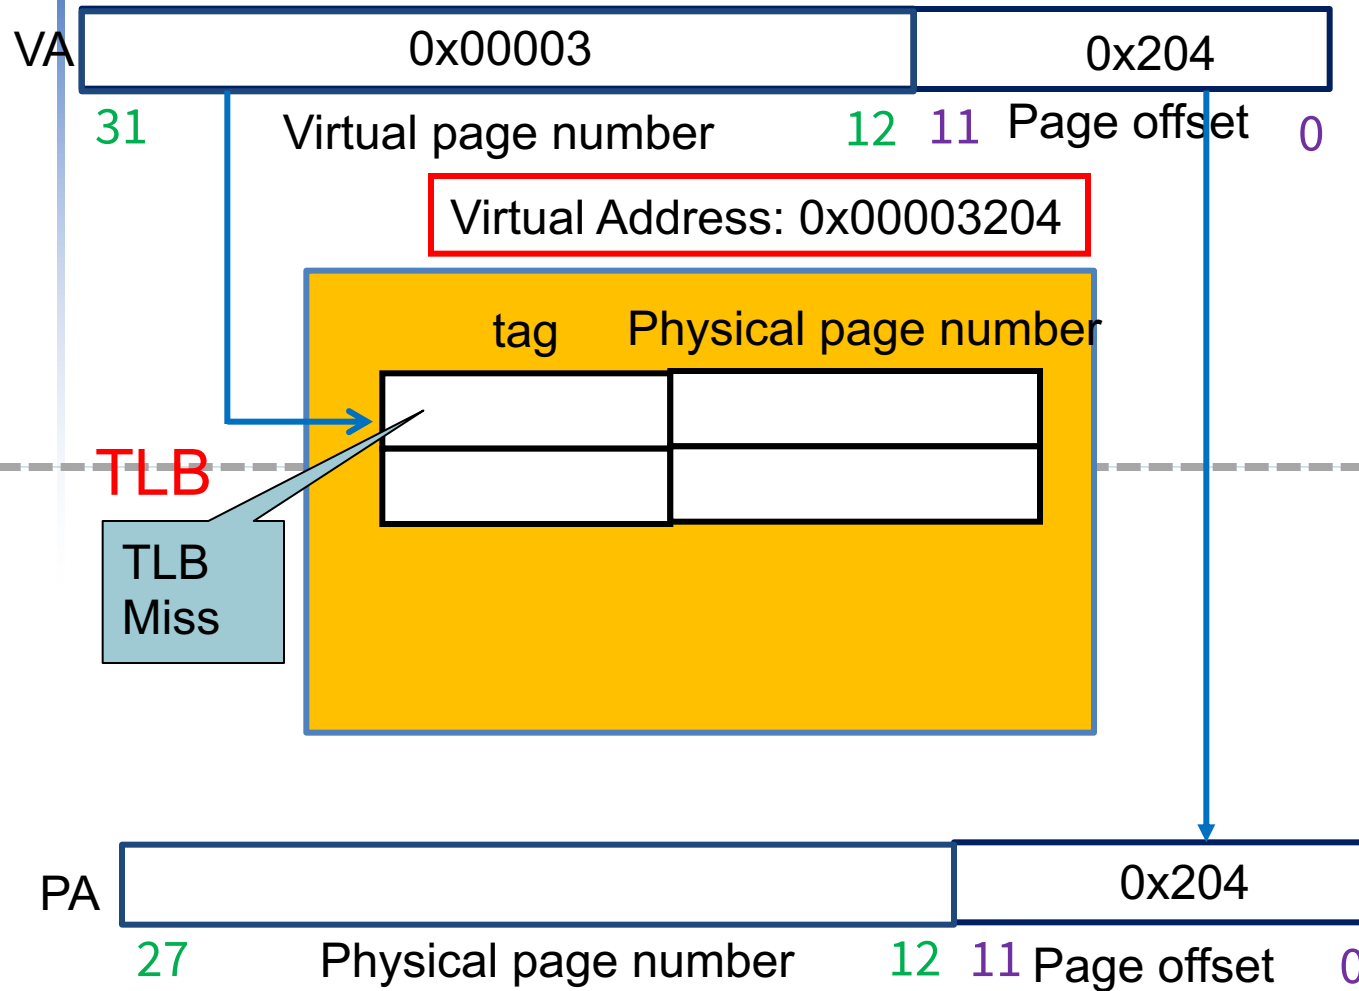
TLB Miss

- Translation not in TLB
 - If page is in memory (page table valid bit = 1)
 - Load the translation from page table to TLB
 - Why not just use page table since it's accessed anyway?
 - Could be handled in hardware or in software
 - If page is not in memory (page fault)
 - Page fault exception – OS handles fetching the page and updating the page table and TLB
 - Then restart the faulting instruction

TLB Miss

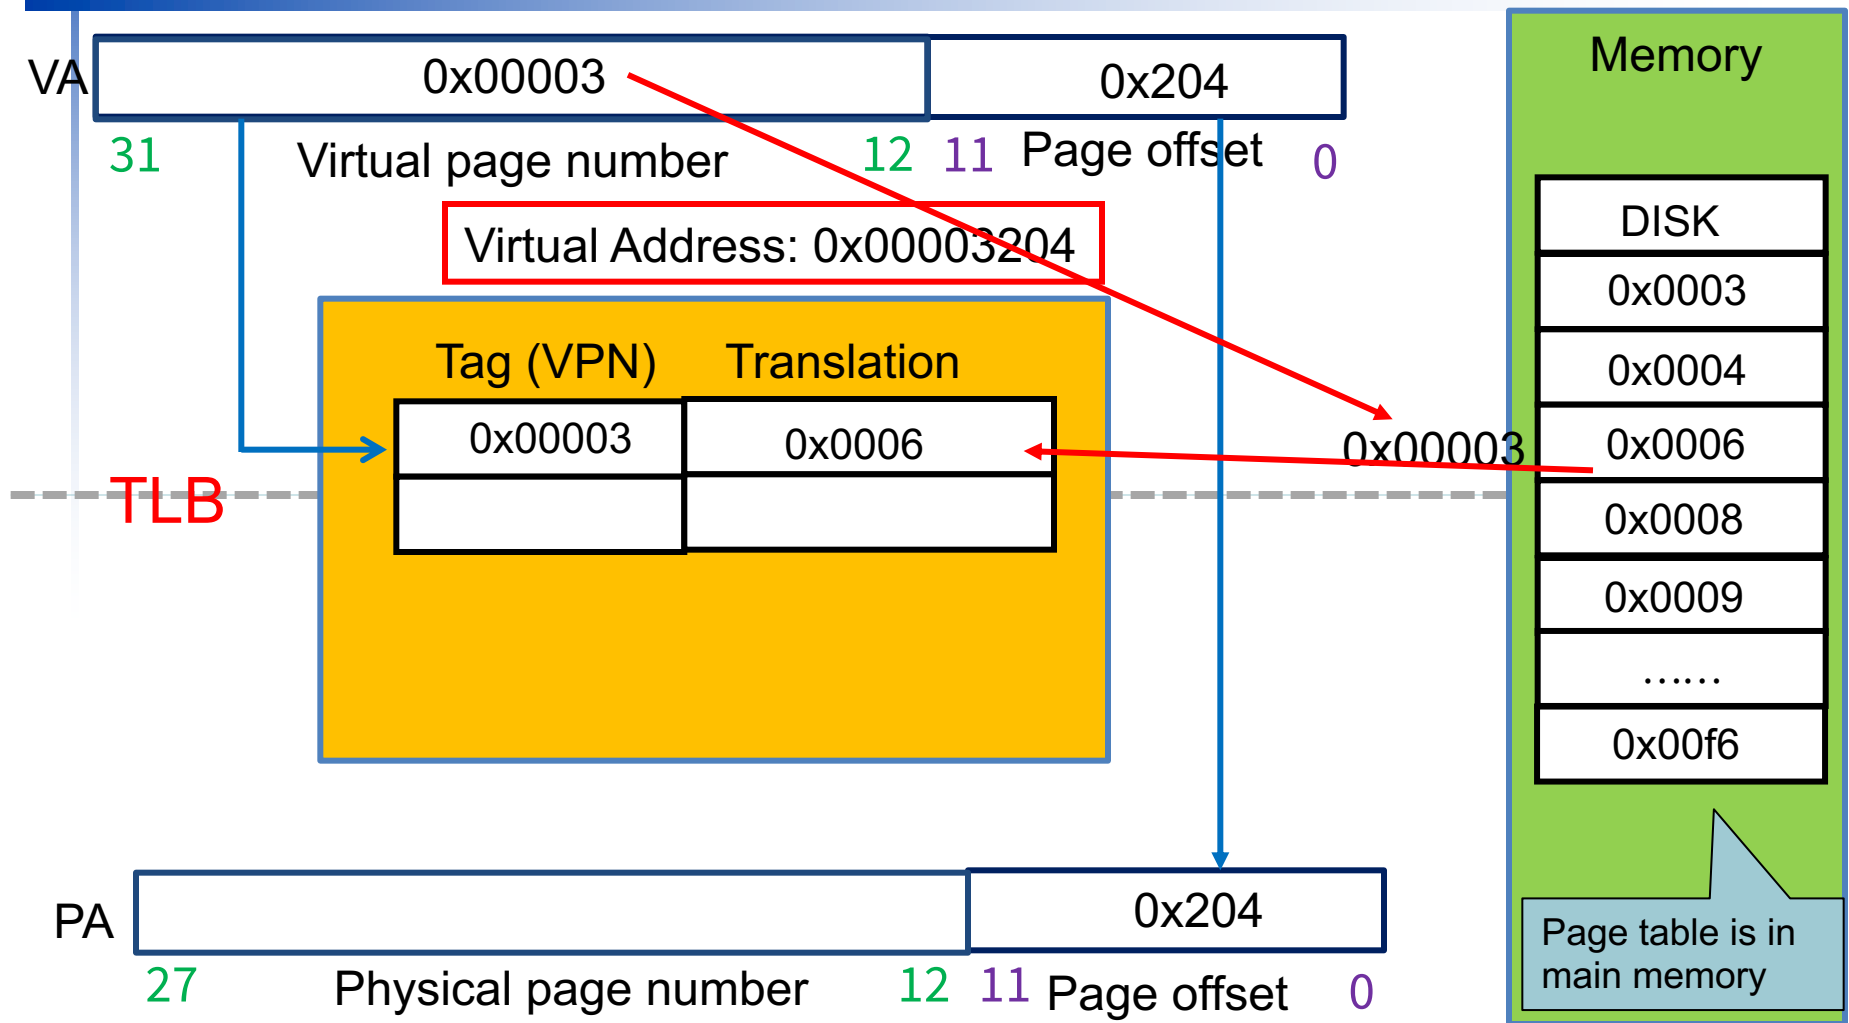
- Recognize a TLB miss or page fault
 - TLB miss – by TLB tag and TLB valid bit
 - Page fault – by page table valid bit

TLB Example (assume 32 bit address)



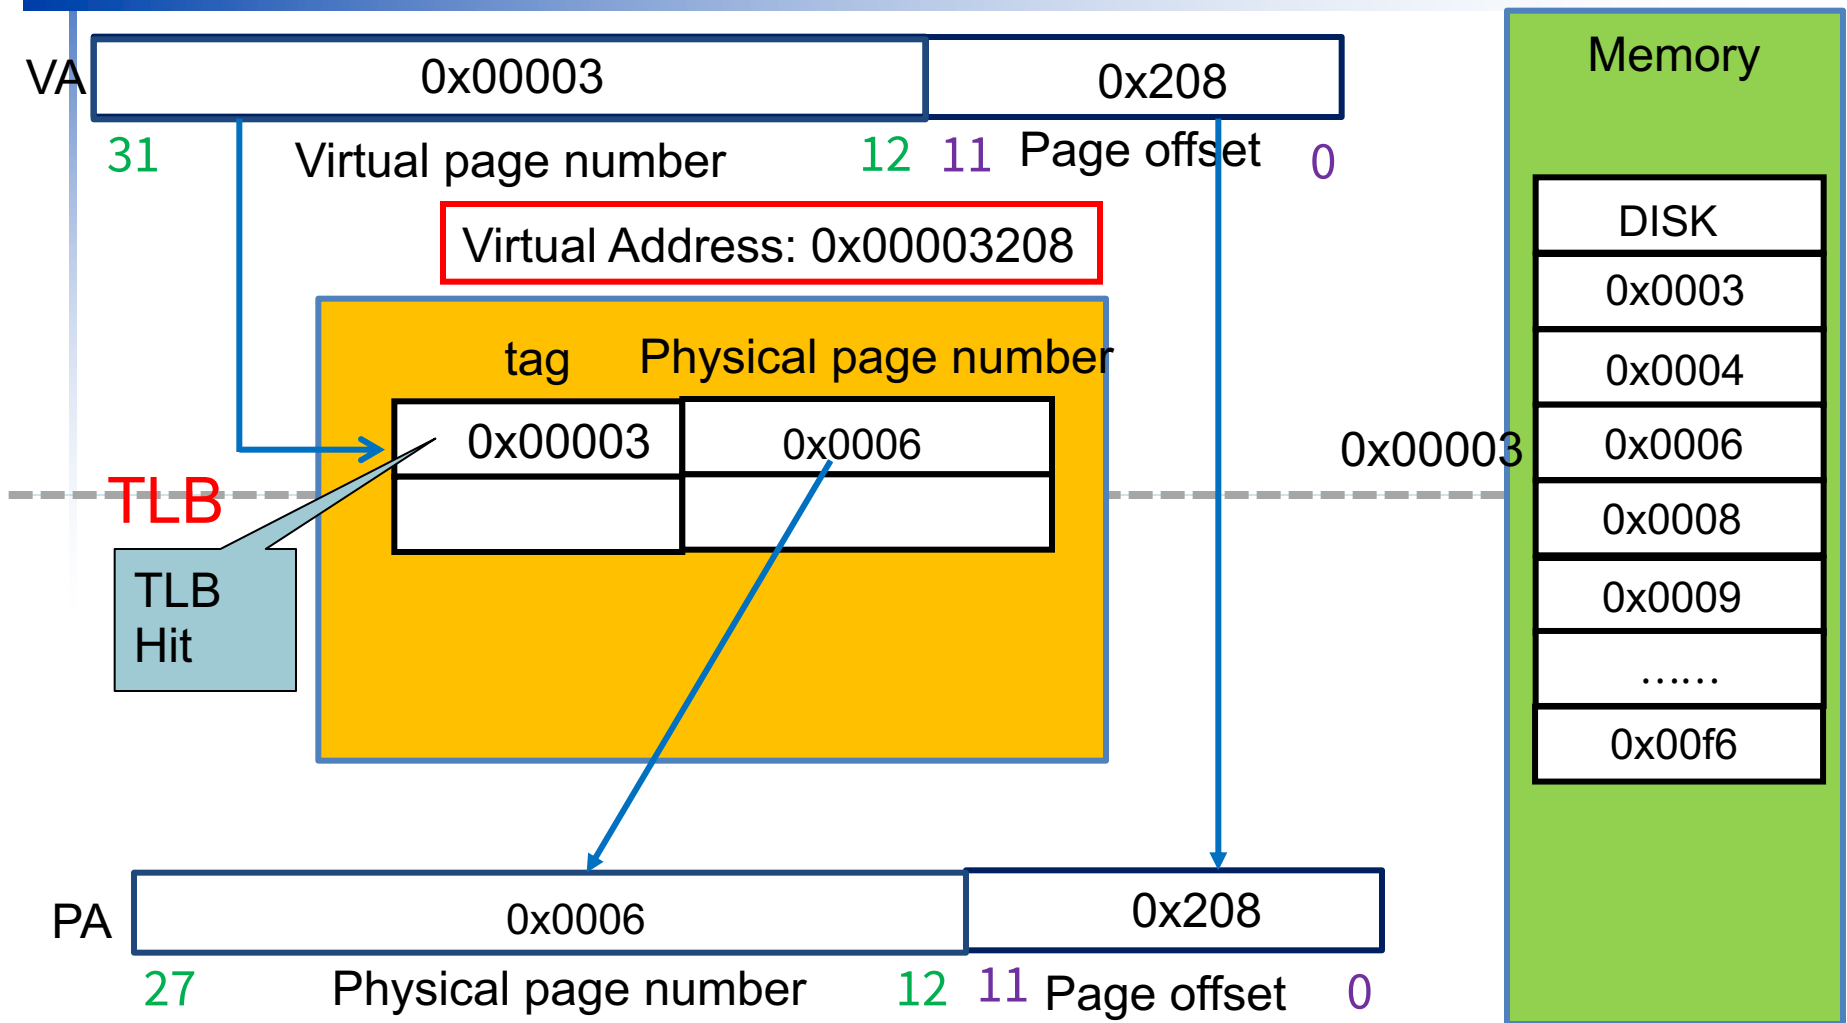
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



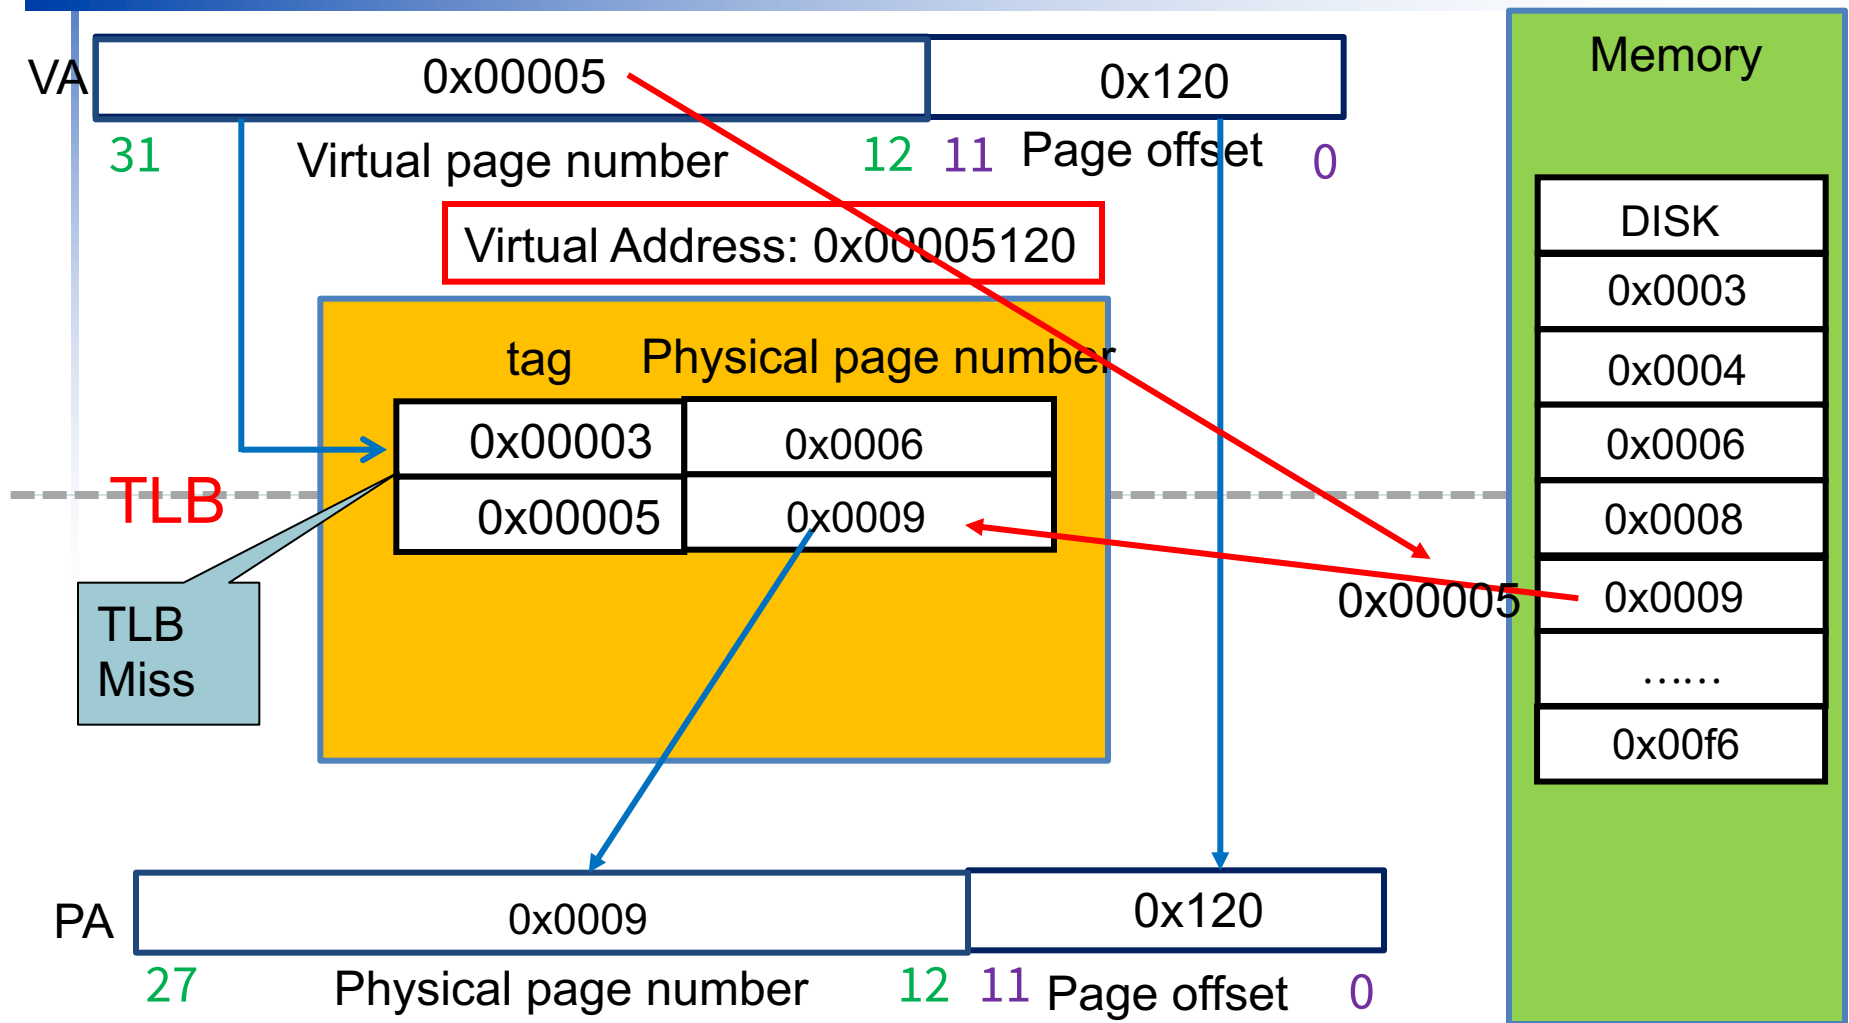
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



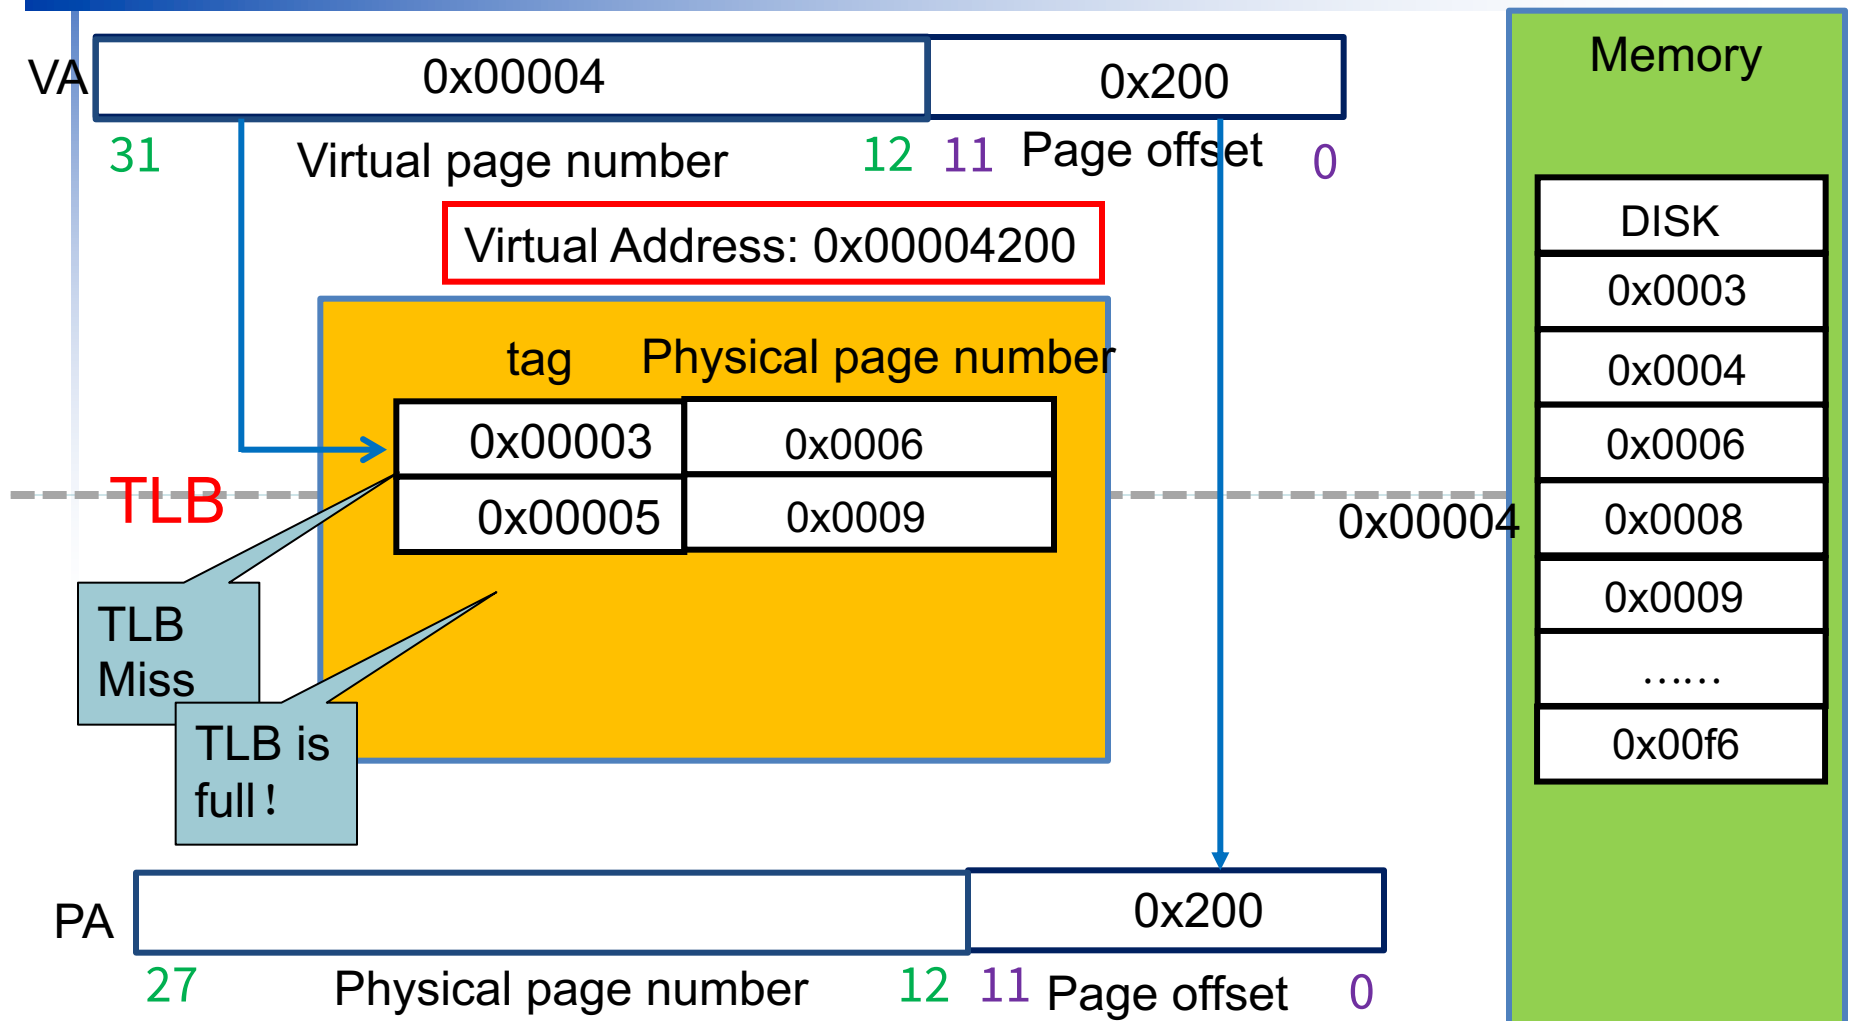
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



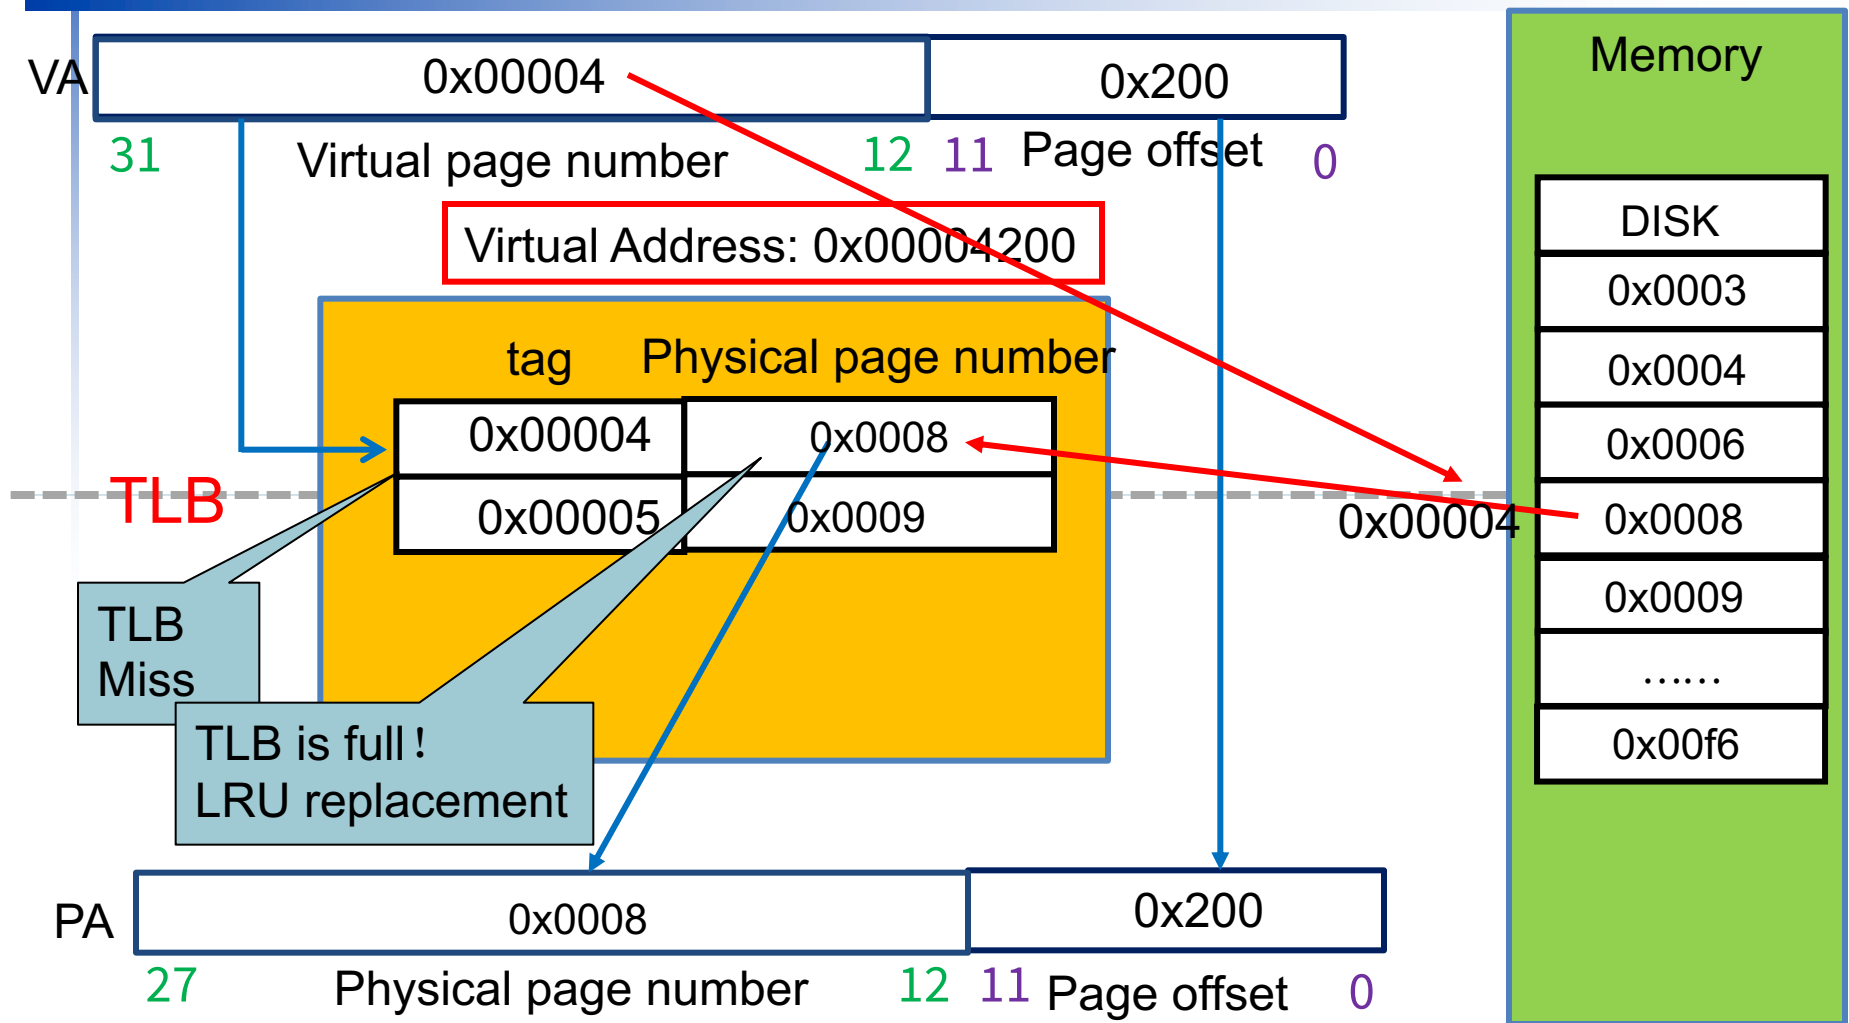
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



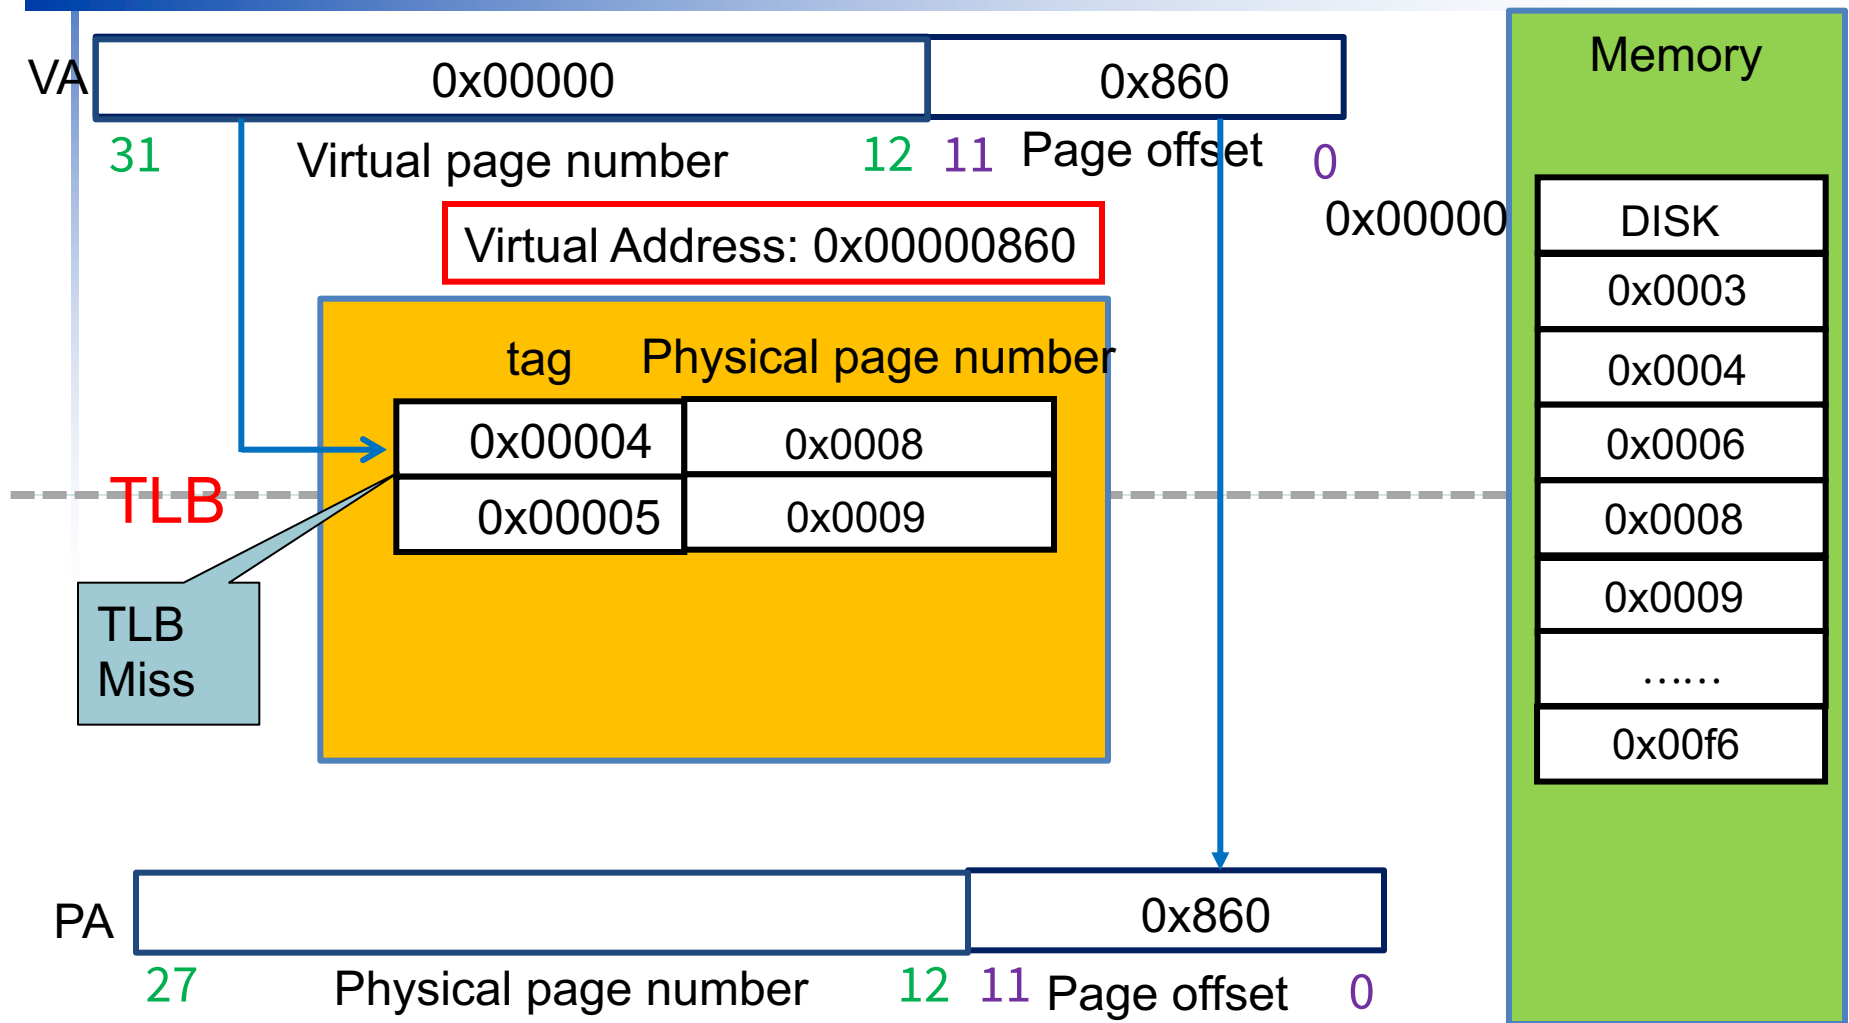
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



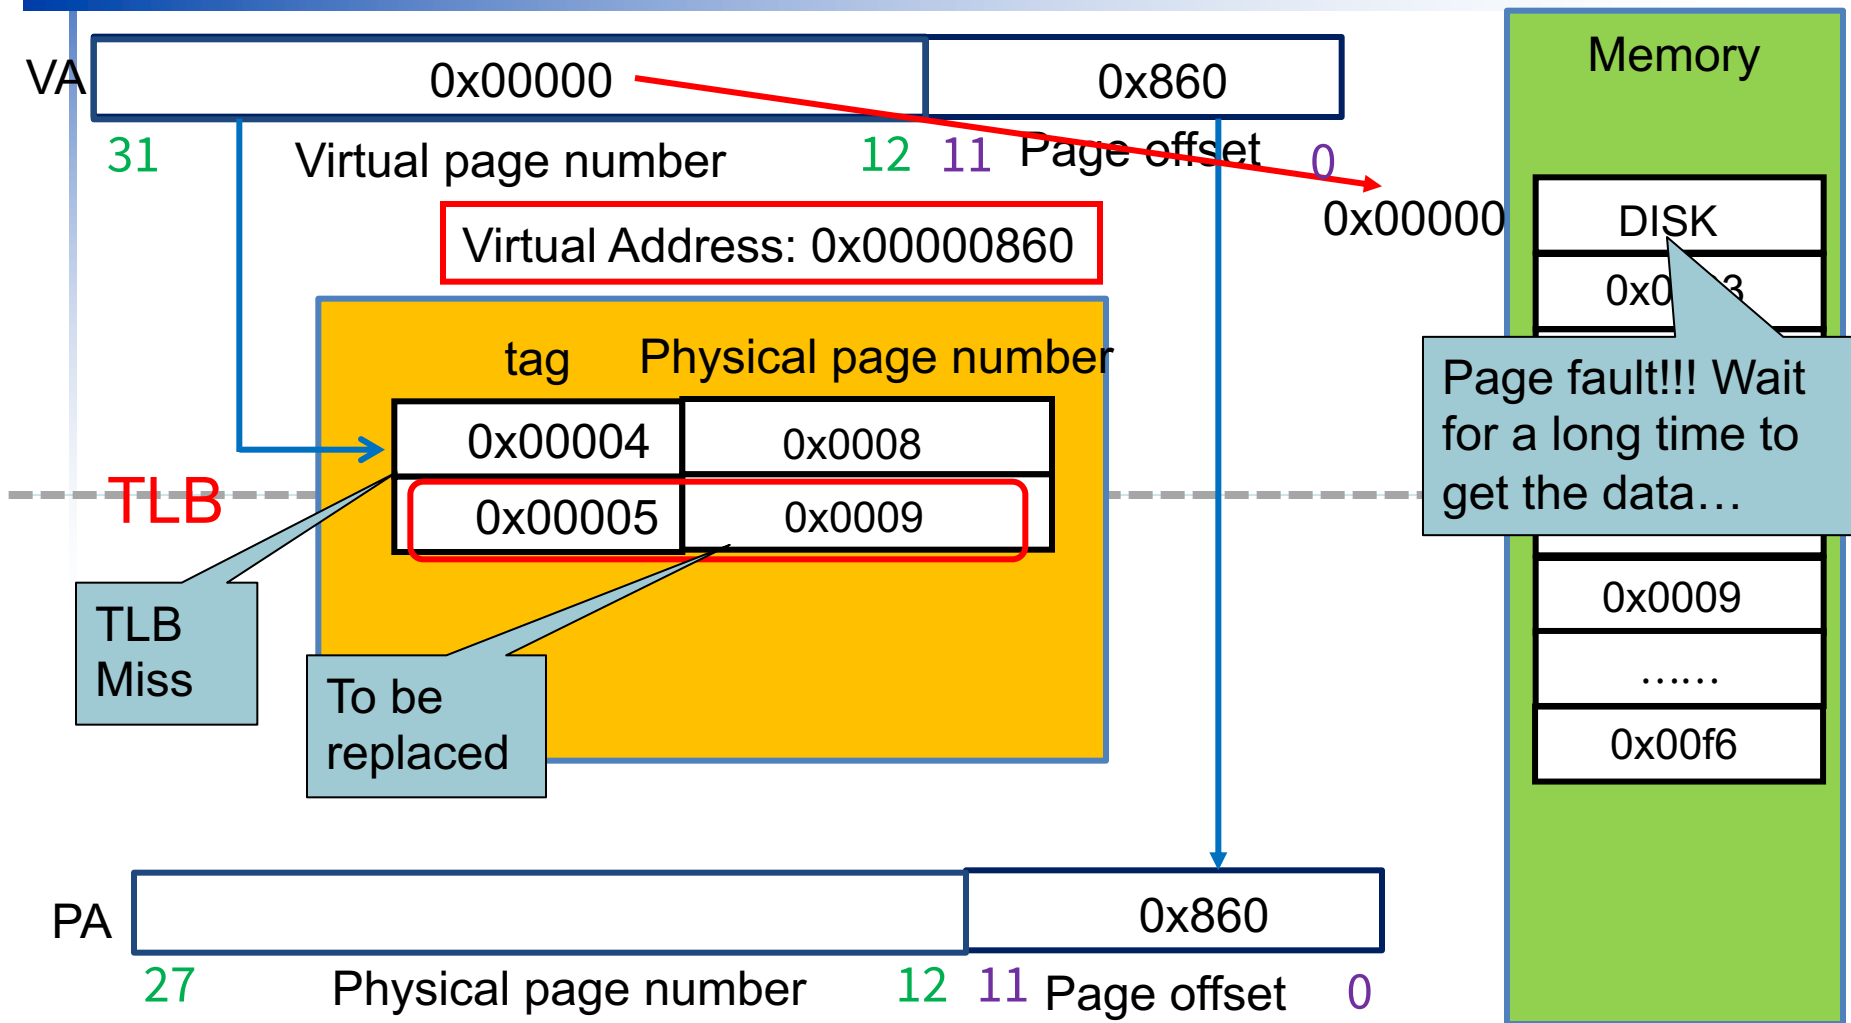
*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)



*Valid, dirty, reference bits omitted here.

TLB Example (cont'd)

- In reality
 - Need to check valid bit, dirty bit and reference bits before overwriting a block on TLB.
- Page fault is unfortunate but unavoidable, it takes a long time to fetch from the disk



Now, we have learned cache, main memory, virtual memory, page table, TLB, ... How do they all work together?



© CanStockPhoto.com

TLB and Cache Interaction

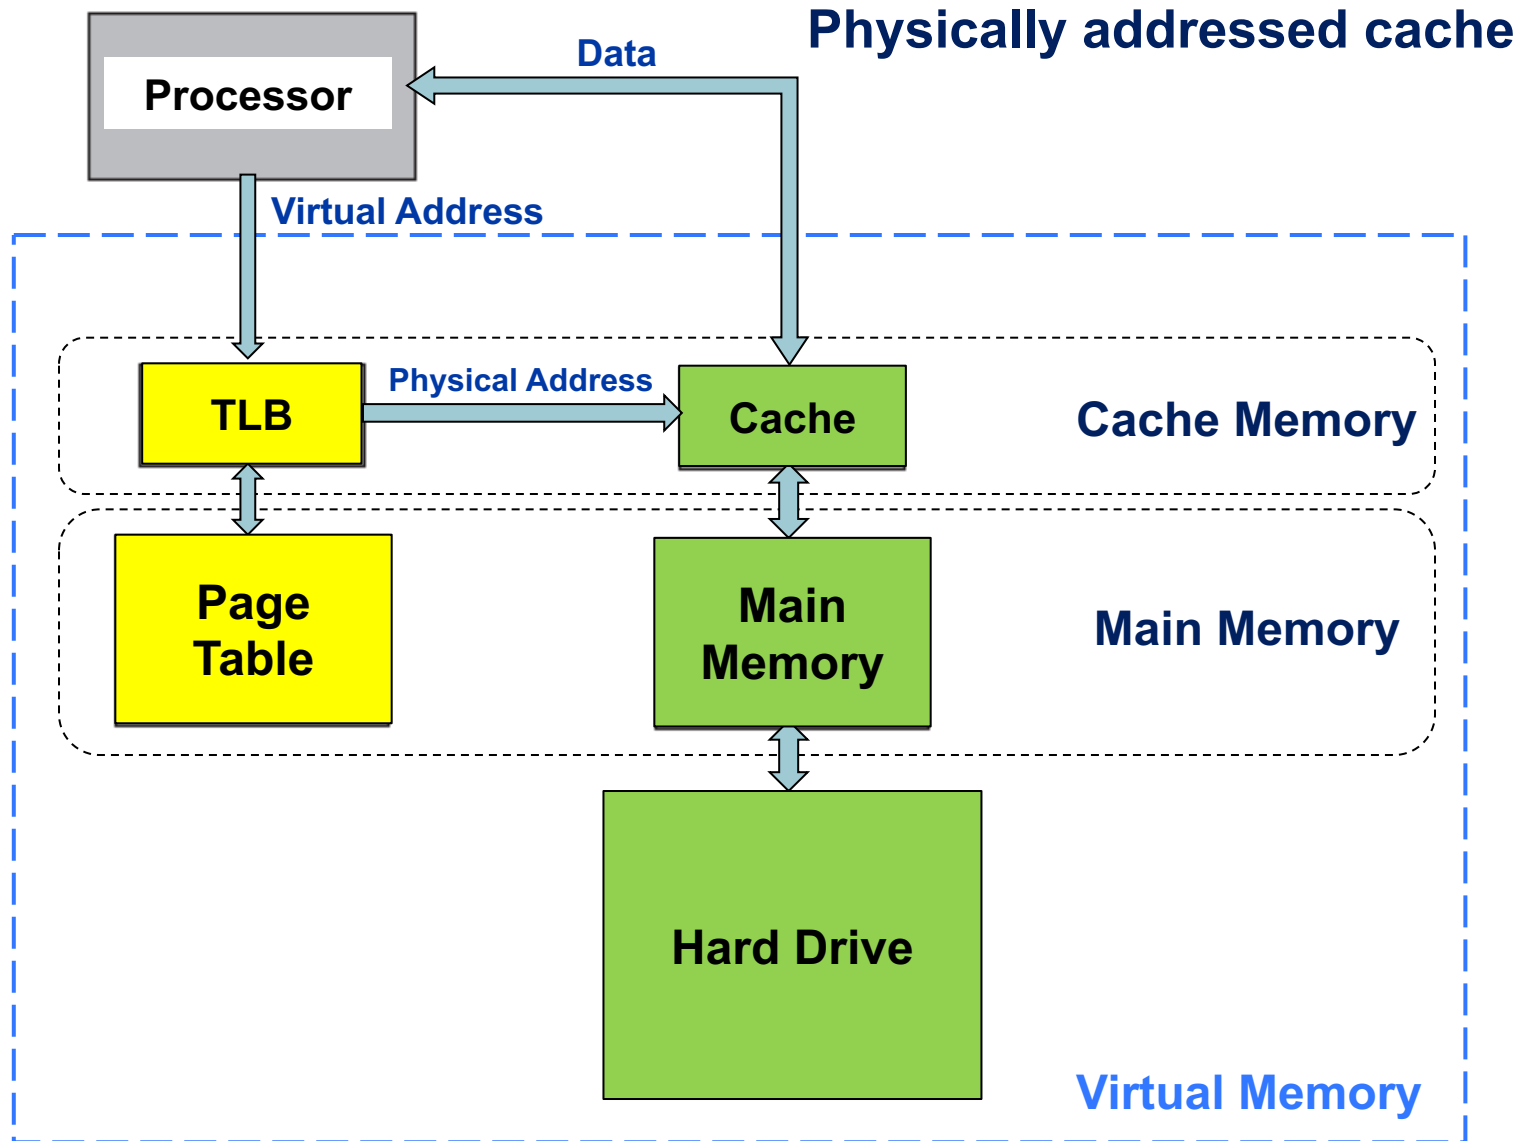
Assuming 4K page size



Small RAM for valid and tags,
Big RAM aligned in word for data

- Indexed by cache index and block offset
- **No need for big MUX, improved cache access time**

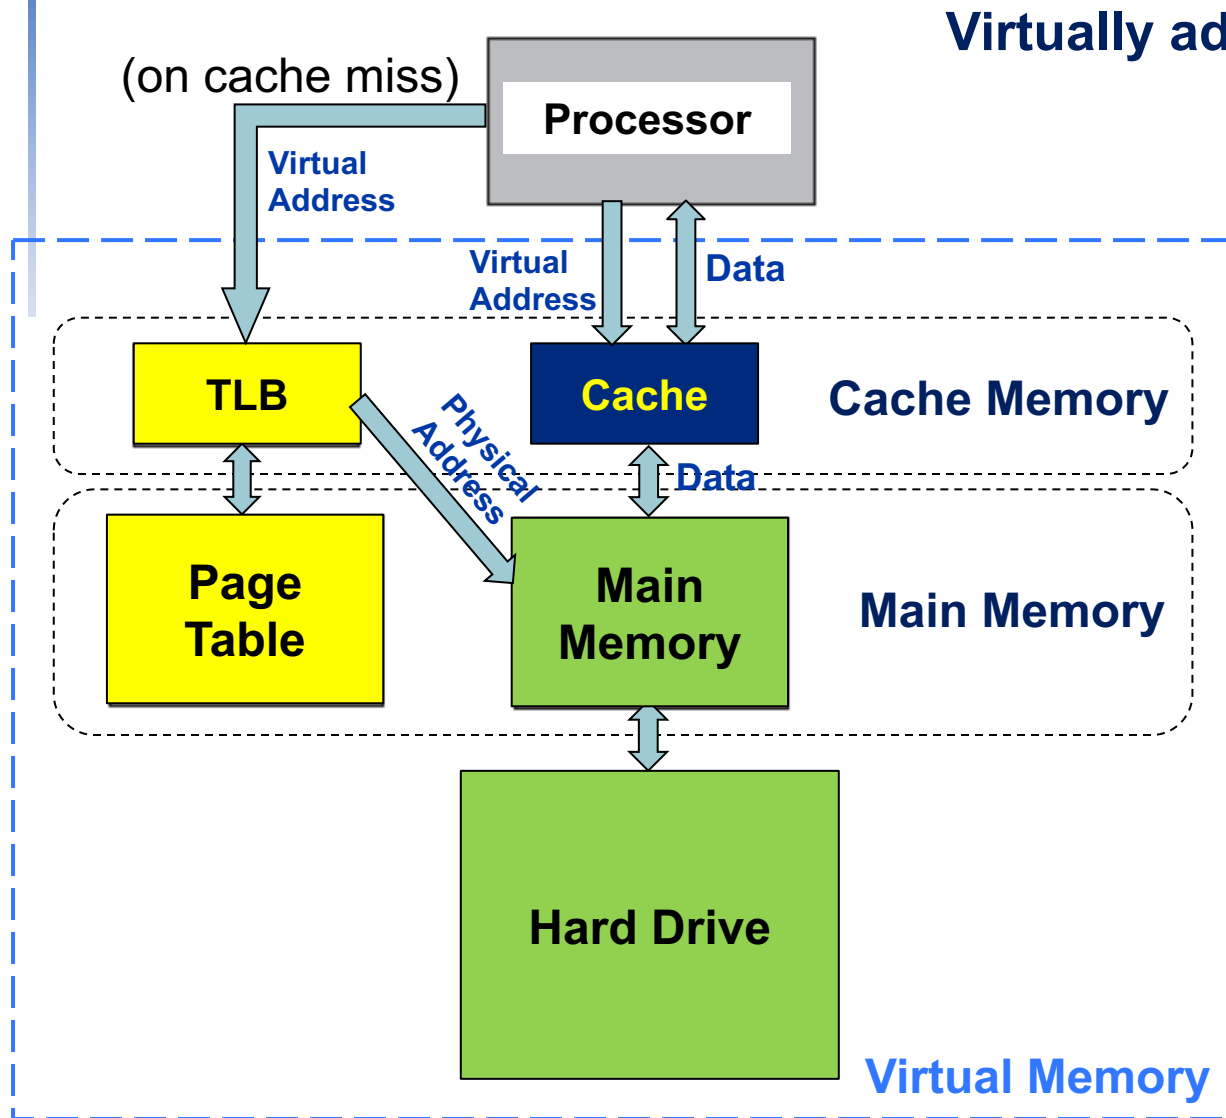
TLB vs. Cache



TLB and Cache Interaction

- If cache uses physical address
 - Need to translate before access cache
 - Virtual Addr → TLB → cache
 - Having TLB on the critical path, taking longer time
- TLB hit and a cache hit are **independent** events, but a cache hit can only occur after a TLB hit occurs, which means that the data must be present in memory.



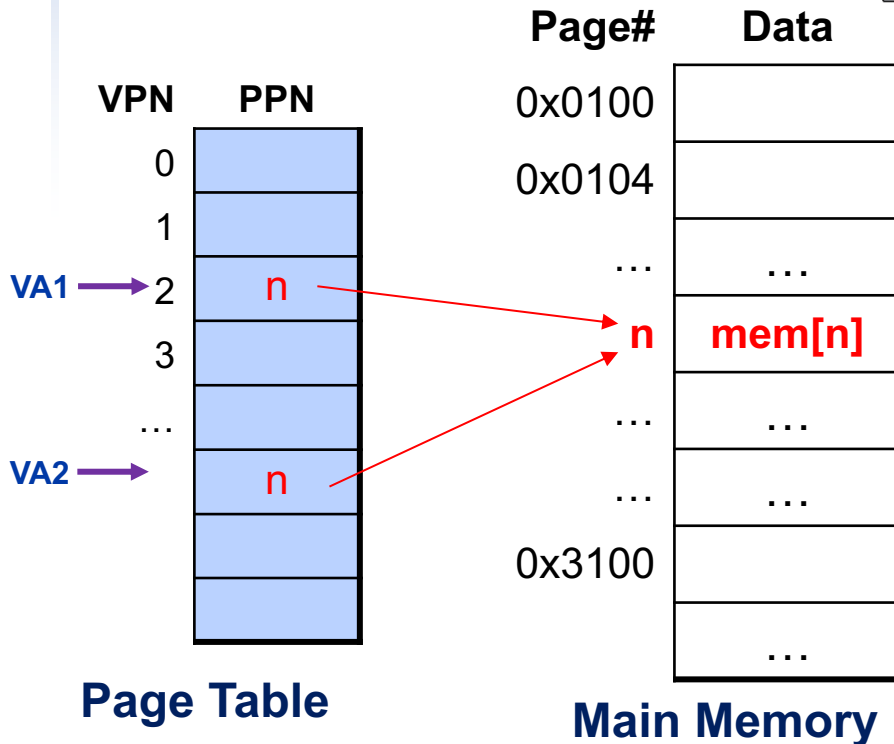
TLB vs. Cache



- No need for translation
- Translate when cache miss
- Complications due to aliasing
 - When different virtual spaces map to the same physical main memory location
- Compromise of both
 - Virtually addressed physically tagged cache

Virtually Addressed Cache

Aliasing: the same object has two different names



Virtually Addressed Cache

Class Exercise with TLB

- Given
 - 4KB page, 16KB physical memory, 4-word blocks, LRU replacement
 - Virtual address: byte addressable, 20 bits (how many bytes?)
 - Page table for program A stored in page #0 of physical memory, starting at address 0x0100, assume only 2 valid entries in page table:
 - Virtual page number 0 => physical page number 1
 - Virtual page number 1 => physical page number 2
 - Fully associative TLB, 2 entries; 4-way associative cache,
- Show how the physical address is used to access a 512-byte cache
- Show the memory structure, complete following table

Virtual Address	Virtual page number	TLB miss?	Page fault?	Physical Address	Cache Hit?	Cache Set Index
0x00F0C						
0x01F0C						
0x20F0C						
0x00100						
0x00200						
0x30000						
0x01FFF						
0x00200						

Relationships in Memory Hierarchy

TLB	Page table	Cache	Possible? If so, under what circumstance?
Hit	Hit	Miss	Possible, although the page table is never really checked if TLB hits.
Miss	Hit	Hit	TLB misses, but entry found in page table; after retry, data is found in cache.
Miss	Hit	Miss	TLB misses, but entry found in page table; after retry, data misses in cache.
Miss	Miss	Miss	TLB misses and is followed by a page fault; after retry, data must miss in cache.
Hit	Miss	Miss	Impossible: cannot have a translation in TLB if page is not present in memory.
Hit	Miss	Hit	Impossible: cannot have a translation in TLB if page is not present in memory.
Miss	Miss	Hit	Impossible: data cannot be allowed in cache if the page is not in memory.

- Assume cache uses physical addresses
- All data in a memory must also be present in its lower level
 - Impossible to copy data cross levels

Memory Protection

- Different programs can have the same virtual addresses
 - But need to protect against errant access
 - Requires OS assistance
- Hardware support for OS protection
 - Privileged supervisor mode (aka kernel mode)
 - Privileged instructions available in supervisor mode
 - System call exception (e.g., `ecall` and `sret` in RISC-V)
 - Page tables and other status information only updated in supervisor mode by operating system
 - Write enable bit protects memory from being written
 - Different page tables ensure different translations
 - Different translation ensures separate memory locations

Summary of The Memory Hierarchy

The BIG Picture

- Common principles apply at all levels of the memory hierarchy
 - Based on notions of caching
- At each level in the hierarchy
 - Block placement
 - Finding a block
 - Replacement on a miss
 - Write policy

Block Placement

- Determined by associativity
 - Direct mapped (1-way associative)
 - One choice for placement
 - n-way set associative
 - n choices within a set
 - Fully associative
 - Any location
- Higher associativity reduces miss rate
 - Increases complexity, cost, and access time

Finding a Block

Associativity	Location method	Tag comparisons
Direct mapped	Index	1
n-way set associative	Set index, then search entries within the set	n
Fully associative	Search all entries	#entries
	Full lookup table	0

- Hardware costs
 - Reduce comparisons to reduce cost
- Virtual memory
 - Lookup table full associativity
 - Benefit in reduced miss rate

Replacement

- Choice of entry to replace on a miss
 - Least recently used (LRU)
 - Needs tracking mechanism (reference bit updated and periodically refreshed)
 - Complex and costly hardware for high associativity
 - Random
 - Close to LRU for high associativity, easier to implement

Write Policy

- Write-through
 - Update both upper and lower levels
 - Simplifies replacement, but may require write buffer
- Write-back
 - Update upper level only
 - Update lower level when block is replaced
 - Need to keep more state
- Virtual memory
 - Only write-back is feasible, given unaffordable disk write latency

Sources of Misses (3C model)

- Compulsory misses (aka cold start misses)
 - First access to a block
- Capacity misses
 - Due to finite cache size
 - A replaced block is later accessed again
- Conflict misses (aka collision misses)
 - In a non-fully associative cache
 - Due to competition for blocks in a set
 - Would not occur in a fully associative cache of the same total size

Cache Design Trade-offs

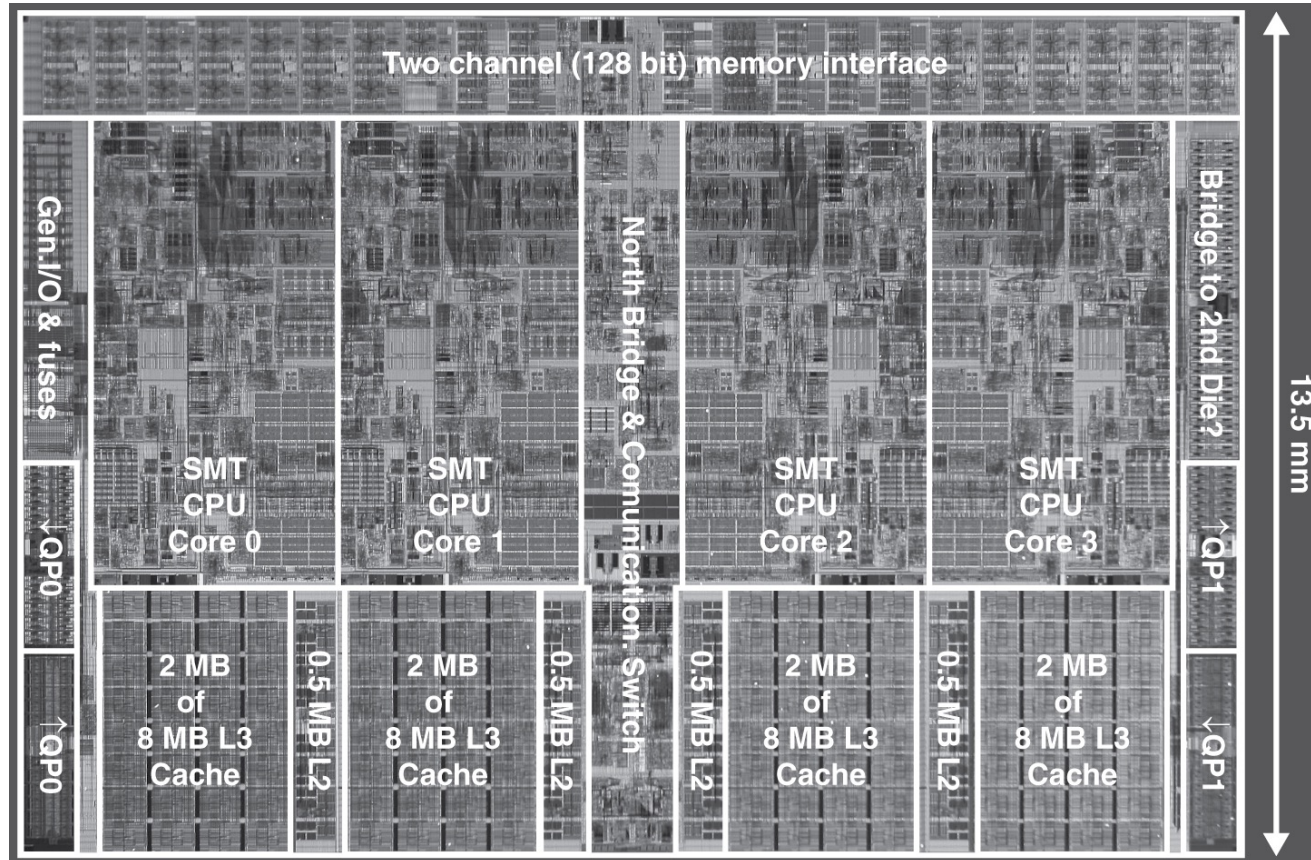
<i>Design change</i>	<i>Effect on miss rate</i>	<i>Negative performance effect</i>
Increase cache size	Decrease capacity misses	May increase access time
Increase associativity	Decrease conflict misses	Will increase access time
Increase block size	Decrease compulsory misses May increase conflict misses	Increases miss penalty. For very large block size, may increase miss rate due to pollution.

Miss Penalty Reduction


- Return requested word first
 - Then back-fill rest of block
- Non-blocking miss processing
 - Hit under miss: allow hits to proceed
 - Miss under miss: allow multiple outstanding misses
 - Need parallel processing capability
- bank interleaved cache/main memory
 - multiple concurrent accesses per cycle

Example: Intel Nehalem 4-core processor

Intel Nehalem 4-core processor



Per core: 32KB L1 I-cache, 32KB L1 D-cache, 512KB L2 cache



2-Level TLB

	Intel Nehalem	AMD Opteron X4
Virtual addr	48 bits	48 bits
Physical addr	44 bits	48 bits
Page size	4KB, 2/4MB	4KB, 2/4MB
L1 TLB (per core)	L1 I-TLB: 128 entries for small pages, 7 per thread ($2\times$) for large pages L1 D-TLB: 64 entries for small pages, 32 for large pages Both 4-way, LRU replacement	L1 I-TLB: 48 entries L1 D-TLB: 48 entries Both fully associative, LRU replacement
L2 TLB (per core)	Single L2 TLB: 512 entries 4-way, LRU replacement	L2 I-TLB: 512 entries L2 D-TLB: 512 entries Both 4-way, round-robin LRU
TLB misses	Handled in hardware	Handled in hardware

3-Level Cache Organization

	Intel Nehalem	AMD Opteron X4
L1 caches (per core)	L1 I-cache: 32KB, 64-byte blocks, 4-way, approx LRU replacement, hit time n/a L1 D-cache: 32KB, 64-byte blocks, 8-way, approx LRU replacement, write-back/allocate, hit time n/a	L1 I-cache: 32KB, 64-byte blocks, 2-way, LRU replacement, hit time 3 cycles L1 D-cache: 32KB, 64-byte blocks, 2-way, LRU replacement, write-back/allocate, hit time 9 cycles
L2 unified cache (per core)	256KB, 64-byte blocks, 8-way, approx LRU replacement, write-back/allocate, hit time n/a	512KB, 64-byte blocks, 16-way, approx LRU replacement, write-back/allocate, hit time n/a
L3 unified cache (shared)	8MB, 64-byte blocks, 16-way, replacement n/a, write-back/allocate, hit time n/a	2MB, 64-byte blocks, 32-way, replace block shared by fewest cores, write-back/allocate, hit time 32 cycles

n/a: data not available

Concluding Remarks

- Fast memories are small, large memories are slow
 - We really want fast, large memories ☹️
 - Caching gives this illusion 😊
- Principle of locality
 - Programs use a small part of their memory space frequently
- Memory hierarchy
 - L1 cache ↔ L2 cache ↔ ... ↔ DRAM memory ↔ disk
 - Virtual space → L1 TLB → L2 TLB → (page table) → physical space
- Memory system design is critical for multiprocessors